

RWTH Aachen University  
Ruei-Bo Chen  
2022/08/15

# ResNet18

## Motivation




Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.




Figure 2. Residual learning: a building block.

# ResNet18

## Motivation




Figure 2. Residual learning: a building block.

$$H(x) = F(x) + x$$

$$F(x) = H(x) - x$$

We want  $F(x)$  to be 0 to make this block an identity mapping. If the network can learn identity mapping, then no matter how many layers we add, they will not affect the convergence.

Why don't we just learn  $H(x)$ ?

# ResNet18

## Motivation




Figure 2. Residual learning: a building block.

$$H(x) = F(x) + x$$

$$F(x) = H(x) - x$$

Why don't we just learn  $H(x)$ ?

For example:

$$x = 2.9, \text{ if } H(x) = 3.0 \rightarrow F(x) = 0.1$$

$$x = 2.9, \text{ if } H(x) = 3.1 \rightarrow F(x) = 0.2$$

$$\Delta H = \frac{3.1 - 3.0}{3.0} = 3.3\%,$$


$$\Delta F = \frac{0.2 - 0.1}{0.1} = 100\%$$

The larger the variance is, the easier the network learns.


# ResNeXt50

## Motivation


Want to increase the accuracy?



change width



change depth

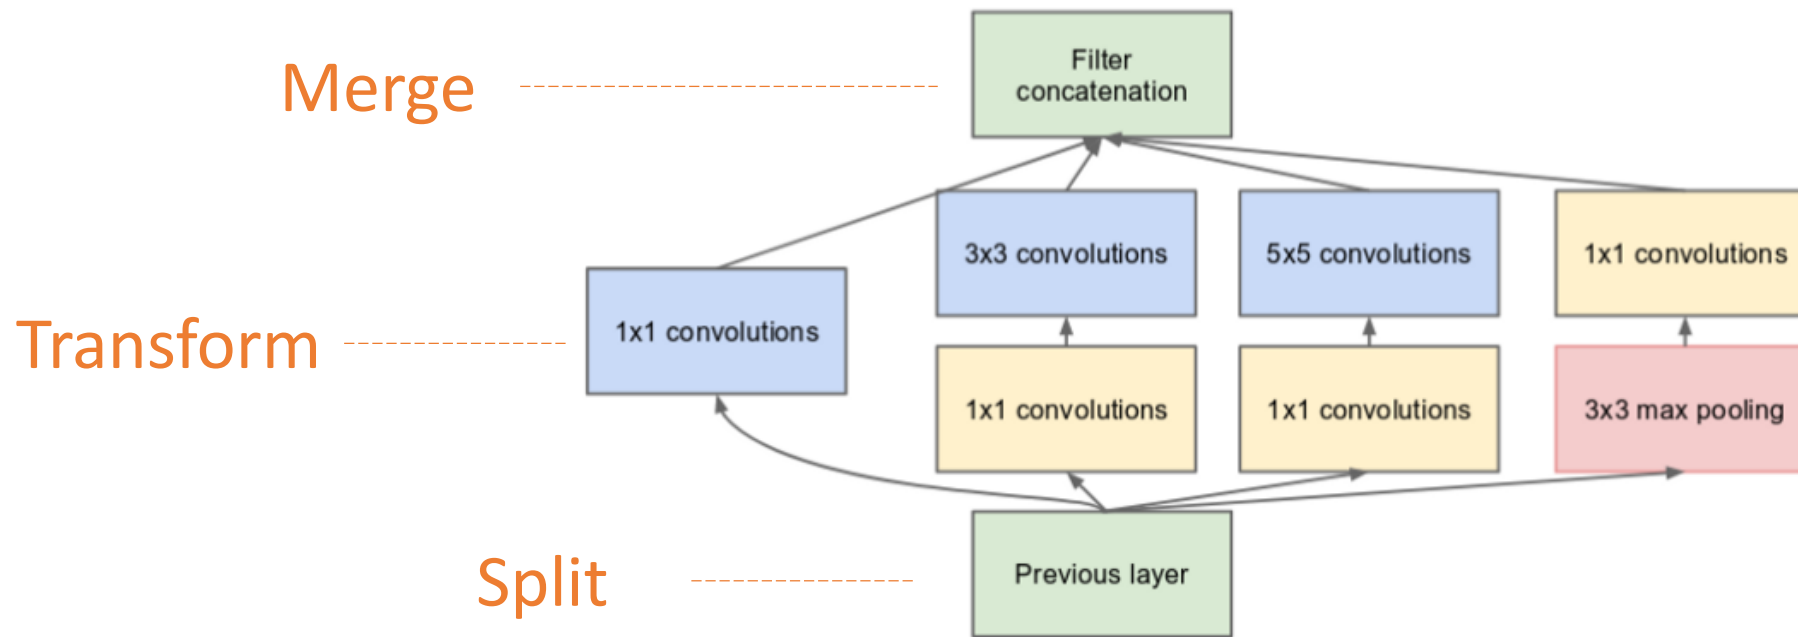


split-transform-merge

# ResNeXt50

## Motivation

### Split-Transform-Merge



(b) Inception module with dimension reductions

# ResNeXt50

## Motivation

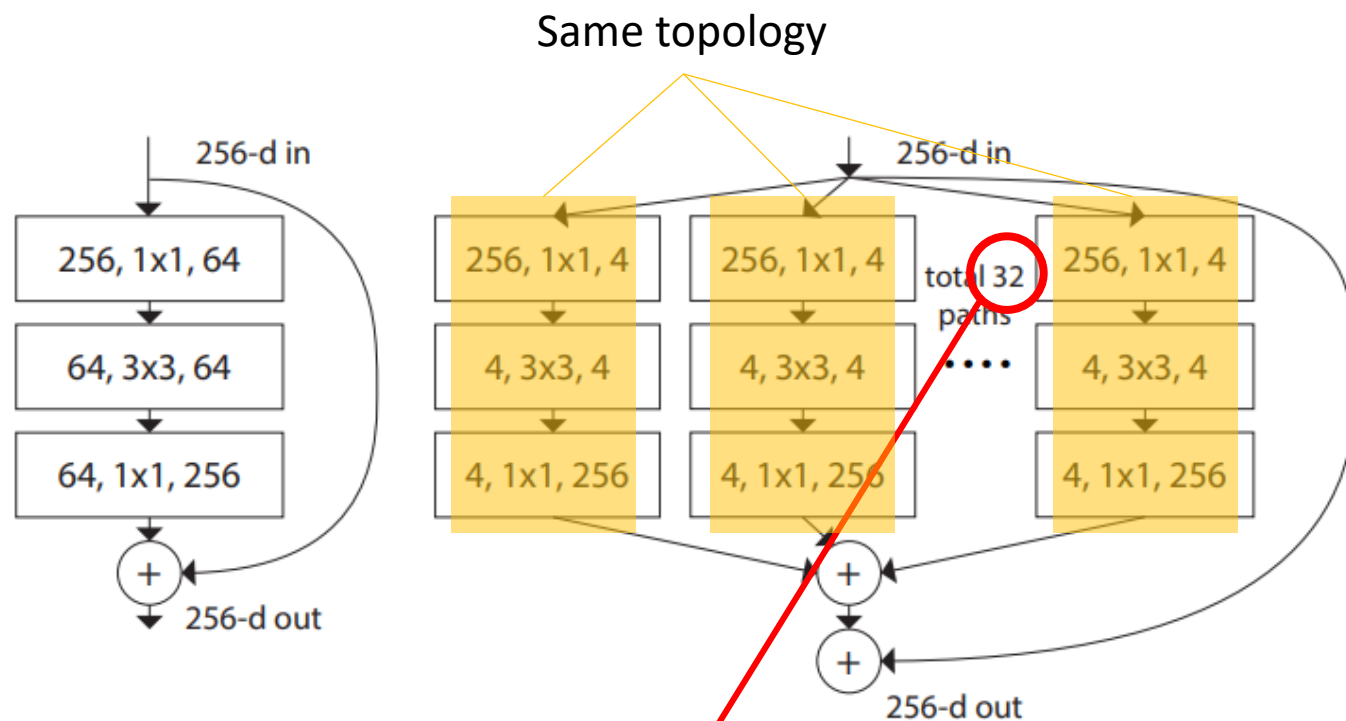


Figure 1. **Left:** A block of ResNet [14]. **Right:** A block of ResNeXt with cardinality = 32 with roughly the same complexity. A layer is shown as (# in channels, filter size, # out channels).

# ResNeXt50

## Motivation




Figure 5. Training curves on ImageNet-1K. **(Left)**: ResNet/ResNeXt-50 with preserved complexity ( $\sim 4.1$  billion FLOPs,  $\sim 25$  million parameters); **(Right)**: ResNet/ResNeXt-101 with preserved complexity ( $\sim 7.8$  billion FLOPs,  $\sim 44$  million parameters).



# ResNeXt50

## Motivation

	setting	top-1 error (%)
ResNet-50	$1 \times 64d$	23.9
ResNeXt-50	$2 \times 40d$	23.0
ResNeXt-50	$4 \times 24d$	22.6
ResNeXt-50	$8 \times 14d$	22.3
ResNeXt-50	$32 \times 4d$	<b>22.2</b>
ResNet-101	$1 \times 64d$	22.0
ResNeXt-101	$2 \times 40d$	21.7
ResNeXt-101	$4 \times 24d$	21.4
ResNeXt-101	$8 \times 14d$	21.3
ResNeXt-101	$32 \times 4d$	<b>21.2</b>

Table 3. Ablation experiments on ImageNet-1K. (**Top**): ResNet-50 with preserved complexity ( $\sim 4.1$  billion FLOPs); (**Bottom**): ResNet-101 with preserved complexity ( $\sim 7.8$  billion FLOPs). The error rate is evaluated on the single crop of  $224 \times 224$  pixels.

	setting	top-1 err (%)	top-5 err (%)
<i>1× complexity references:</i>			
ResNet-101	$1 \times 64d$	22.0	6.0
ResNeXt-101	$32 \times 4d$	21.2	5.6
<i>2× complexity models follow:</i>			
ResNet- <b>200</b> [15]	$1 \times 64d$	21.7	5.8
ResNet-101, wider	$1 \times$ <b>100d</b>	21.3	5.7
ResNeXt-101	<b>2</b> $\times 64d$	20.7	5.5
ResNeXt-101	<b>64</b> $\times 4d$	<b>20.4</b>	<b>5.3</b>

Deeper—  
Wider—

Table 4. Comparisons on ImageNet-1K when the number of FLOPs is increased to  $2 \times$  of ResNet-101's. The error rate is evaluated on the single crop of  $224 \times 224$  pixels. The highlighted factors are the factors that increase complexity.

# ResNeXt50

## Motivation

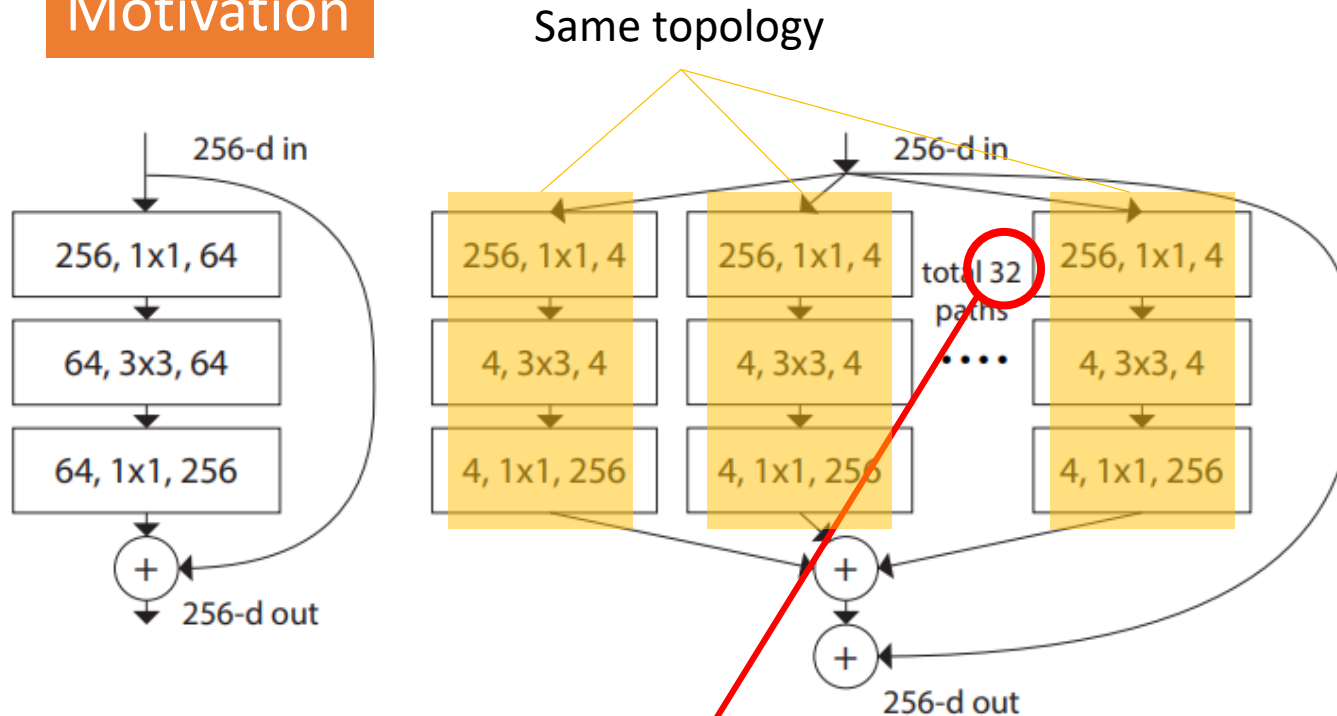


Figure 1. **Left:** A block of ResNet [14]. **Right:** A block of ResNeXt with cardinality = 32 with roughly the same complexity. A layer is shown as (# in channels, filter size, # out channels).

1. Combines the greatness of Inception and ResNet. (Allow the model to jointly attend to the information from different feature subspaces)
2. Because the topology of each cardinality is the same, the model reduce the number of hyperparameters, make it more general in other case.
3. Increase the accuracy without increasing the number of parameter.

# Comparison

```
model_name = 'resnext'  
continued_model = initialize_model(model_name, num_classes, use_pretrained=False)  
continued_model.load_state_dict(torch.load('./models/resnext_best/best_weight.pt').state_dict())  
continued_model = continued_model.to(device)  
model_testing(continued_model, data_loader)
```

Test complete in 0m 42s  
Acc: 0.7071

```
model_name = 'resnet'  
test_model = initialize_model(model_name, num_classes, use_pretrained=False)  
test_model.load_state_dict(torch.load('./models/res18_best/res18best_weight.pt').state_dict())  
test_model = test_model.to(device)  
model_testing(test_model, data_loader)
```


Test complete in 0m 19s  
Acc: 0.7162

In terms of the test result, ResNet18 performs better and faster!

1. ResNet18 is faster because its complexity is smaller than ResNeXt50 naturally.
2. ResNeXt is not trained perfectly. The training overfits in the end. It could still be optimized.

# Improve the performance

## 1. L2-Regularization for overfitting



ResNeXt Overfitting

```
scratch_optimizer = optim.Adam(scratch_model.parameters(), lr=0.001)
```

Add 'weight\_decay=5e-4' as an argument

# Improve the performance

## 2. Add more data or use stronger and variant data augmentation

E.g., SnapMix

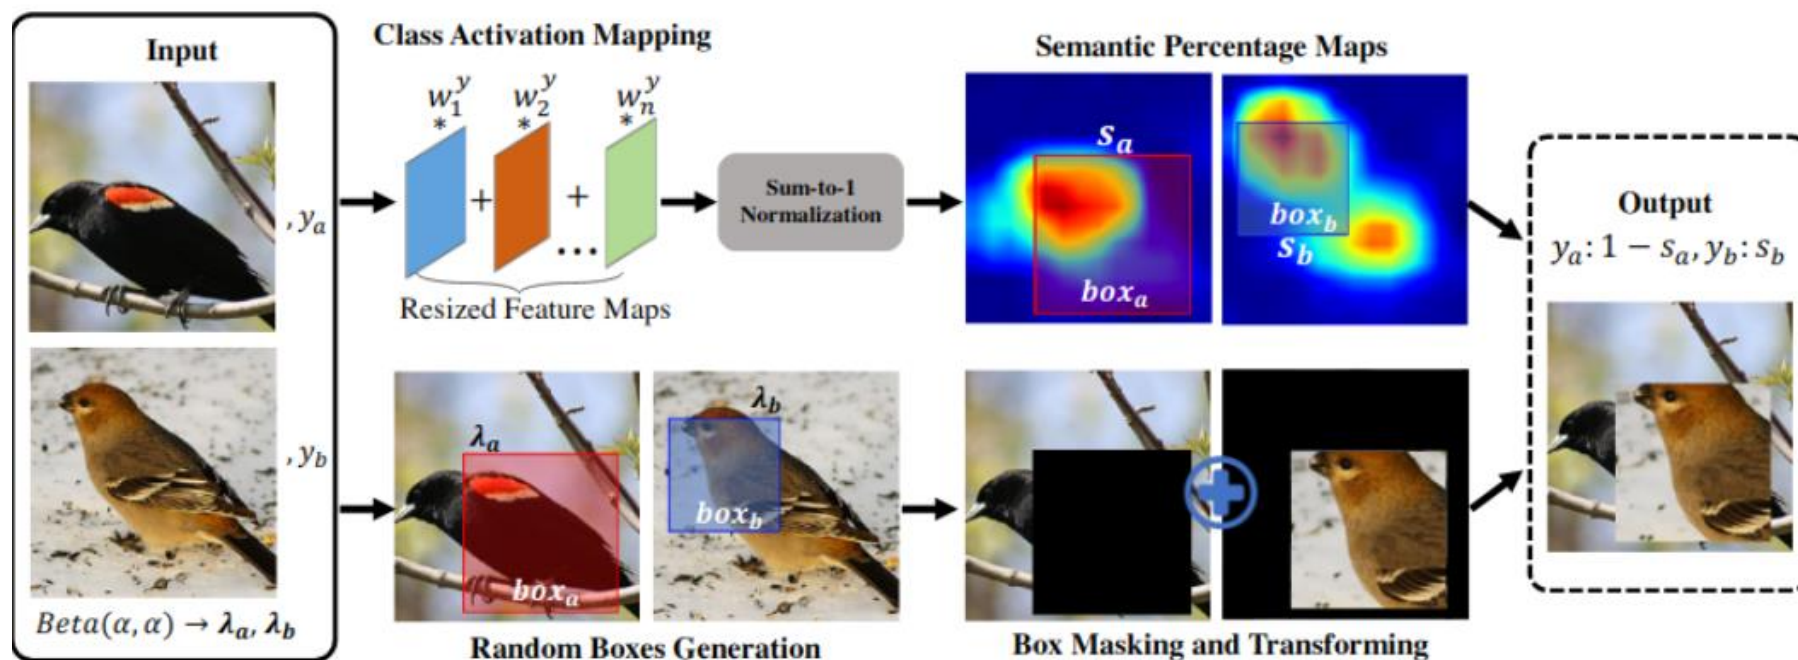


Figure 2: An overview of proposed method.

# Improve the performance

## 3. AutoML for better hyperparameter



Cloud AutoML Vision

AutoML can assist developers to tune the hyperparameter in a better way.

The most naïve way is to test all of the possible hyperparameters automatically.

We can try to use AutoML to tune ResNeXt. Since its number of hyperparameter is not that large, it would be a good way.

# Improve the inference speed

## 1. GPU Utilization

NVIDIA-SMI 450.51.06      Driver Version: 450.51.06      CUDA Version: 11.0									
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC			
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.		
0	Tesla V100-SXM2...	On	00000000:8A:00.0	Off	62%	0		Default	
N/A	49C	P0	142W / 300W	2880MiB / 16160MiB				N/A	
Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
	ID	ID				Usage			
0	N/A	N/A	64709	C	python	2873MiB			

GPU utilization will dramatically affect the inference speed. Bad transfer of data from the host to device can result in low GPU utilization.

1. Utilize prefetch technique to make CPU preprocess the data while waiting for GPU training current batch.
2. Cache those data which don't need to be preprocessed in each epoch in the memory. (Tensorflow has this option, not sure whether Pytorch has a similar function?)
3. Use bottleneck to check which part of the data pipeline is stuck for the longest time.

# Improve the inference speed

## 2. Mixed Precision

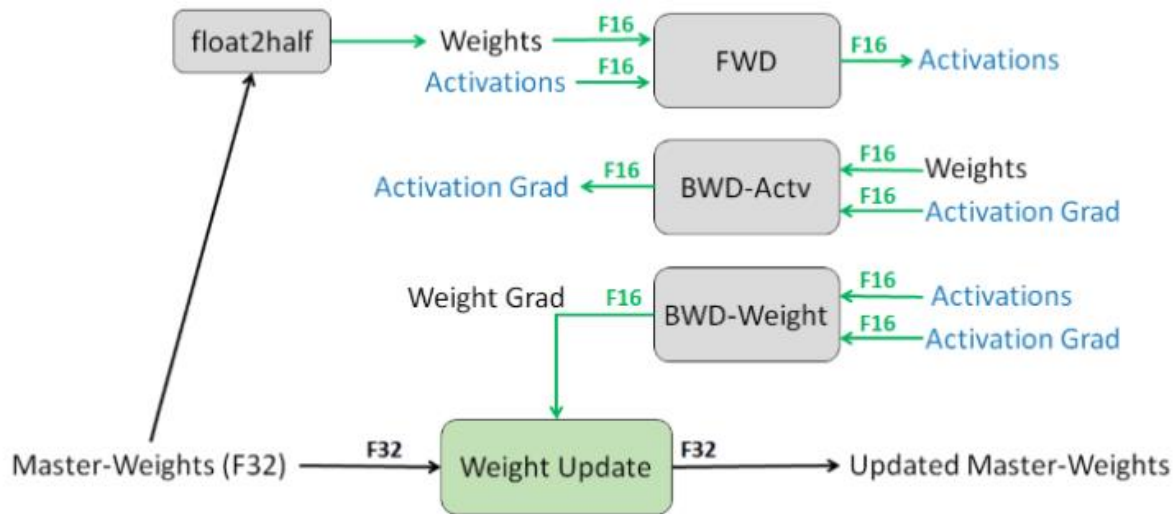


Figure 1: Mixed precision training iteration for a layer.

Use Float16 to replace Float32 during training, saving half of the space and improve the throughput.

1. Pytorch has the option to implement mixed precision conveniently.
2. When enabling mixed precision, remember to adjust the output type of the model layers.